



Artificial Neural Networks with small Datasets. A practical Approach

Masterarbeit

zur Erlangung des akademischen Grades
Master of Science in Engineering (M.Sc.)

Eingereicht bei:

Fachhochschule Kufstein Tirol Bildungs GmbH
Data Science & Intelligent Analytics

Verfasser:

Paul Leitner, BA

1910837299

Erstgutachter : Dr. Johannes Luethi
Zweitgutachter : Lukas Demetz, PhD

Abgabedatum:

31. October 2021

Eidesstattliche Erklärung

Ich erkläre hiermit, dass ich die vorliegende Masterarbeit selbstständig und ohne fremde Hilfe verfasst und in der Bearbeitung und Abfassung keine anderen als die angegebenen Quellen oder Hilfsmittel benutzt sowie wörtliche und sinngemäße Zitate als solche gekennzeichnet habe. Die vorliegende Masterarbeit wurde noch nicht anderweitig für Prüfungszwecke vorgelegt.

Kufstein, 31. October 2021

Paul Leitner, BA

Contents

1	Introduction & Problem Statement	1
1.1	Objectives & Methods	4
2	Data augmentation and Data synthesis in other Areas	5
2.1	Synthetic Data for Privacy / Data Protection	6
2.1.1	Synthetic Data Vault	8
2.1.2	DataSynthesizer	9
2.1.3	Efficacy of synthesized data	11
2.2	Synthetic Minority Oversampling Technique	12
2.3	Data Augmentation for Image Data	14
3	Generative Adversarial Networks	17
3.1	A brief summary of Generative Adversarial Networks	18
3.1.1	The Generator Network	18
3.1.2	The Discriminator Network	18
3.1.3	GAN loss function	19

3.1.4	GAN training loop	20
4	Technical Application	22
4.1	Theoretical applicability	22
4.2	Technical implementation steps	23
4.2.1	Network Architecture	24
4.2.2	Network Training Implementation	32
5	Data Boosting Experiment Results	40
5.1	Experiments on the original dataset	40
5.1.1	Experiments with decreased amounts of data	41
5.2	Different Model Types and Datasets	51
5.2.1	Reference - model fitting on the titanic dataset	55
5.3	Fitting the Generator and models on completely new data	57
5.3.1	The Wine Quality Dataset	57
5.3.2	The Pima Indians Diabetes Dataset	61
5.4	Replacing Training Data with Synthetic Data	65
5.4.1	Models on purely synthetic Titanic Data	66
5.4.2	Models on purely synthetic Wine Data	67
5.4.3	Models on purely synthetic Diabetes Data	68
5.4.4	Replacement Conclusions	69

6 Discussion	72
6.1 General Thoughts	72
6.2 Conclusions from architecture and technical design	74
6.2.1 Gridsearch	77
6.3 Conclusions from shrinking and boosting the original dataset .	78
6.4 Conclusions from application to other datasets	80
6.5 Conclusions from complete data replacement	83
6.6 Final Thoughts	85

FH Kufstein Tirol

Data Science & Intelligent Analytics

Abstract of the thesis: **Artificial Neural Networks with small Datasets. A practical Approach**

Author: Paul Leitner, BA

First reviewer: Dr. Johannes Luethi

Second reviewer: Lukas Demetz, PhD

After giving a summary on the literature and history of other approaches for increasing training data, I elucidate the trade-offs between deep learning and other machine learning approaches. I show that advanced machine learning approaches such as Neural Networks ([NNs](#)) mostly trade increased data requirements against data scientist effort in data preparation and feature engineering on the one hand and model complexity and representational power on the other hand.

I then investigate whether more complicated [NNs](#) may be used by synthetically enlarging the training data present using Generative Adversarial Networks ([GANs](#)) and thereby achieving comparable or greater accuracy while saving data preparation time, effectively trading processing time (synthetic data enlargement being resource-intensive) for manual feature-engineering time by creating a model and benchmarking it against shallow reference models on multiple standard Machine Learning ([ML](#)) datasets.

31. October 2021

1. Introduction & Problem Statement

In their paper "The unreasonable effectiveness of Data" Pereira et al. (2009) (alluding to "The unreasonable effectiveness of Mathematics in the Natural Sciences" by Wigner (1990)) the authors posed one of the fundamental axioms in machine learning today, (often even invoked as shorthand) the fact that a larger set of training data would **always** be preferable, all other things being equal.

Ultimately, predictive power of machine learning models seems to be driven by

- Size and quality of training data
- Expressive power of models (itself limited by processing capabilities)

That is, assuming that a model which fits the task at hand exists and has sufficient power to learn, the deciding factor in successful application of machine learning becomes the quality and amount of data.

The number of books (such as the one by Brownlee (2020)) on the topic of data preparation - that is, reformulation of existing data to facilitate the extraction of patterns from it - gives some indication on the importance of this step.

The example given by Géron (2019) in his book of predicting time of day from the image of a clock face illustrates this most clearly: while it would be possible to train a neural network on a large number of images of clock faces, forcing them to infer the times the hands of the clock point at eventually - by simply reformulating the image into a polar coordinate system of hands pointing at certain angles, however, the problem may be solved directly with a few lines of code.

Munson (2012) showed in their survey on the importance of different steps in the process of creating a predictive model, that the preparation of data occupies a significant share of time spent by data scientists attempting to apply machine learning methodologies.

Of significance is the implied notion that while data preparation is a comparatively time-intensive part of the process of machine learning, this is the case since effort spent in data preparation is a lever to increase model performance through more data scientist work. Not all steps of the process offer the same (potential) tradeoff between work and predictive power. In the same paper Munson (2012) notes that model tuning only consumes on the order of 14% of data scientist's time. This allocation of time follows from the assumption that it is **more likely** that the ultimate degree of success on a task will be affected by enriching the data at hand than it is to innovate an entirely new method of prediction, data clustering or inference.

Furthermore, while most books on the topic mention the acquisition of additional data and the benefits of enriching and reformulating the data at hand, it is almost axiomatic that the objective is making ideal use of the data at hand.

The amount of data that is necessary to solve a given task can usually not be defined, Raudys et al. (1991) attempt to quantify the effect of small sample size, on classification tasks specifically.

Data size and model complexity are usually inversely correlated depending on the task. Scikit-learn, the de-facto standard library for python "shallow" machine learning (that is, not using NNs) offers a "decision map" on the right model:

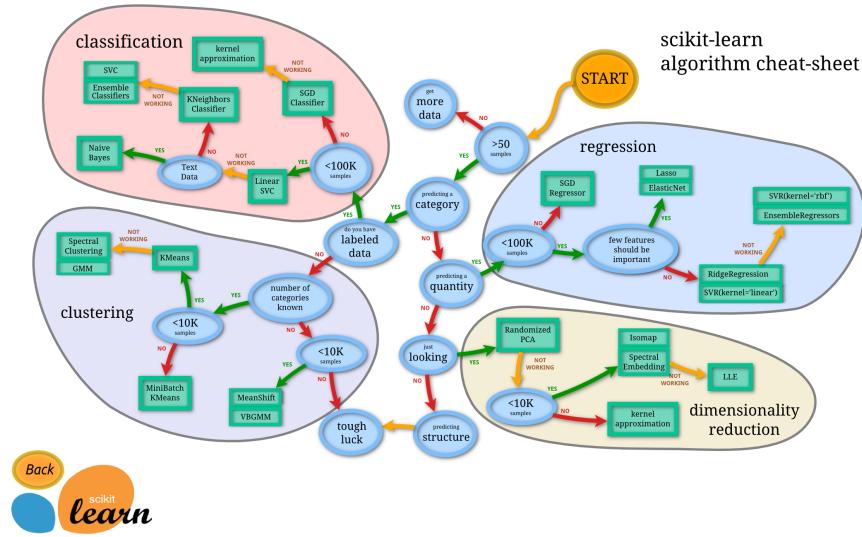


Figure 1: Decision Map for choosing the right estimator, from scikit learn
https://scikit-learn.org/stable/tutorial/machine_learning_map/index.html

The takeaway here is the fact that many key "decision nodes" here are based on the data size available.

The goal of this paper is to empirically test whether or not this strict relationship may be ameliorated by boosting the available training data.

1.1 Objectives & Methods

In contrast to "shallow" learning methods, deep learning refers to creating intermediate representations in **NNs** via **stacked layers** at the cost of increased training data as **Swingler (1996)** alludes to.

In his excellent book on deep learning, keras author **Chollet (2017)** construes neural networks as structured hierarchical models which extract progressive representations from data, thereby enabling a more abstract understanding of patterns within the data.

The objective of this paper is therefore to test whether more complex deep learning models may be applied to a small dataset by synthetically enlarging the training data, effectively bypassing the need for extensive feature engineering by employing an outsized amount of computational power to a small dataset, effectively trading computation for data preparation time.

After a review of different approaches to creating synthetic data and their results in a later chapter (2), the practical viability of this approach will be tested practically in chapter 5.

The method used to increase training data is a **GAN** - to be introduced in a later chapter (3). The **GAN** will be developed on a well-known small machine learning dataset. After boosting the dataset size, different models will be trained on this dataset to compare their performance. If the results are promising, the method will be encapsulated into a python package and published on the Python Packaging Index (**pypi**) - <https://pypi.org/>.

2. Data augmentation and Data synthesis in other Areas

Due to the importance of the amount of training data available and its' large impact on machine learning outcomes, using existing training data to maximum effect is a topic that has been approached from multiple sides. More established approaches involve K-Fold Crossvalidation as explored by [Anguita et al. \(2009\)](#) and bootstrap aggregating ([bagging](#)).

In these approaches, broadly speaking, the available data is sub-sampled in different ways to simulate the presence of multiple - subtly differing - training sets, which makes them analogous to [NN](#) dropout as described by [Srivastava et al. \(2014\)](#) in their paper and ensemble methods as described by [Dietterich \(2000\)](#).

Creating entirely new synthetic training data or **altering** existing data to fit machine learning models is separate area of research.

2.1 Synthetic Data for Privacy / Data Protection

A significant amount of the impetus behind this area of research stems from privacy concerns and attendant legislation. Especially noteworthy here is of course the General Data Protection Regulation ([GDPR](#)) which imposes potentially devastating penalties on the misuse of data regarding to natural persons ("natural person" in the legal sense, as opposed to a legal person such as a corporation).

A full discussion of data privacy is out of scope in this paper, but the motivation to sidestep concerns involving specific persons which arises from the potential sanctions arising from the [GDPR](#) or US Health Insurance Portability and Accountability Act ([HIPAA](#)) and similar legislation by data anonymization is clear and understandable.

[El Emam et al. \(2020\)](#) explores the topic well in his book. As illustrated here:

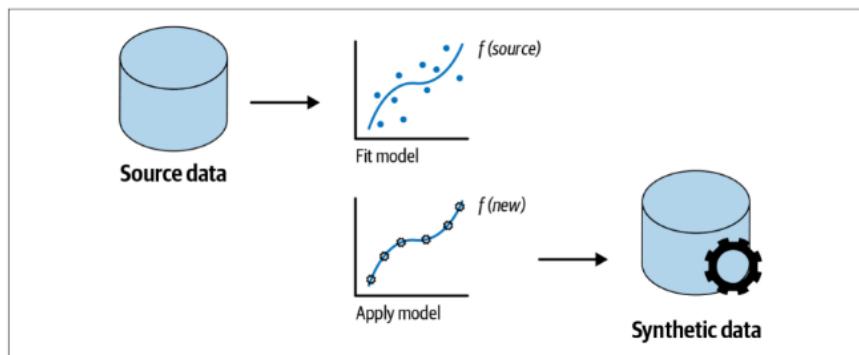


Figure 2: from [El Emam et al. \(2020\)](#)

As the graphic shows quite succinctly, the logical process involves extracting relevant patterns from the original data, explicitly or implicitly, and then producing new data which preserves these patterns. The generated data may then be used, without legal or ethical concerns regarding the privacy interests of the subjects of the original data, in order to extract insights and conduct research.

The implicit assumption behind this concept is, however: that the patterns which are relevant may be known beforehand and/or separated from the types of research conducted on the new data. More directly put, extracting patterns from a table of data with some attributes may be fairly trivial since there is only a limited if not always small number of interrelations possible between multiple attributes. The situation changes dramatically if the data in question is stored in a relational database or distributed system, or is semi-structured or unstructured.

In these cases, it may be difficult to even identify patterns in the data which need to be preserved in order to facilitate its' use in possible use cases in the future.

This assumption may not always hold true but there is existing research which has been conducted to mitigate this problem.

2.1.1 Synthetic Data Vault

Patki et al. (2016) describe an interesting approach to solve the issue of anonymization of an entire relational database specifically.

Their approach, briefly, proposes to model relations between tables specifically, in an iterative manner:

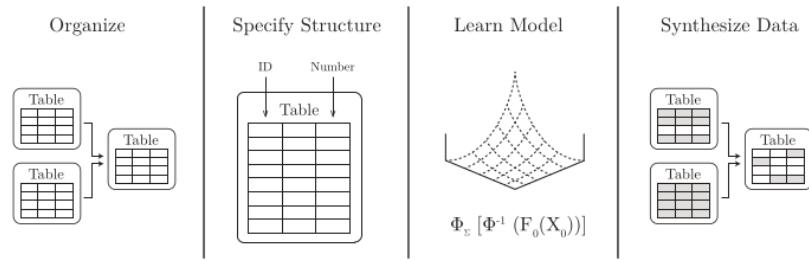


Figure 3: from Patki et al. (2016)

The model of the database then exposes these tables via an Application Programming Interface (API) to be queried and generates synthetic tables based on the tables queried and their interrelations. In their paper Patki et al. (2016) assess the efficacy of their approach by crowd-sourcing prediction tasks on datasets generated with their approach to data scientists and find that in over 50% of cases there was no significant difference in predictive power between features generated on the original dataset vs. generated data (an independent comparison to another approach to synthetic data created for privacy is found here: 2.1.3).

2.1.2 DataSynthesizer

Ping et al. (2017) propose a method with a stronger focus on tabular data in their paper.

Their method, the DataSynthesizer focuses more strongly on the inference of data types, supporting most common types

- integer
- float
- string
- datetime

and then infers their meaning from there, for example attempting to recognize whether an attribute that is stored as a string describes a category.

The resulting generation and assessment process is illustrated here:

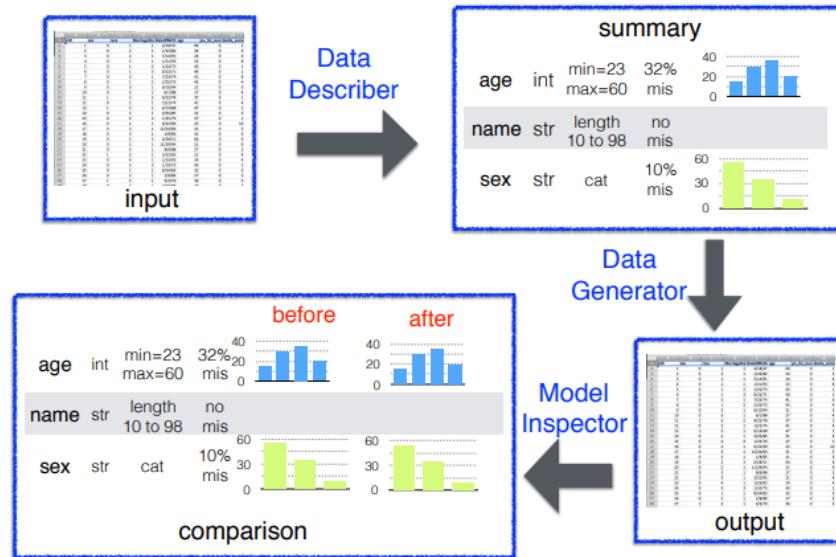


Figure 4: from Ping et al. (2017)

It is really quite an elegant approach which attempts to produce a statistically similar generated dataset.

2.1.3 Efficacy of synthesized data

Hittmeir et al. (2019) compare these two approaches to each other in detail and assess their efficacy in creating classification models in comparison to each other and to real data:

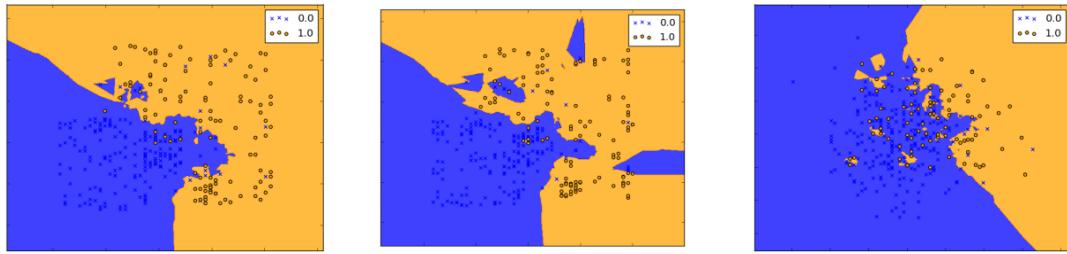


Figure 5: showing the decision boundary of an **SVM** classifier trained on data generated on original data (left), data generated using the **DS** (middle) and **SDV** (right) from Hittmeir et al. (2019)

As a note here, Hittmeir et al. (2019) used tabular datasets for evaluation which plays to the strength of the **DS** approach. Some further notes in the conclusion, their paper is quite pertinent to the objective here.

2.2 Synthetic Minority Oversampling Technique

A popular approach in which synthetic data is used, synthetic minority oversampling technique ([SMOTE](#)) focuses on strongly imbalanced datasets. As [Chawla et al. \(2002\)](#) argue in their paper, real-world data is often imbalanced in such a way that the case which needs to be detected constitutes only an extremely small part of the data, while detection of these samples is critical. Noteworthy examples include the detection of cancer and fraudulent bank transactions.

The focus of [SMOTE](#) is on the identification of interesting feature subspaces, and not only sampling from these regions with small perturbations in the data, while undersampling the majority class but actually creating more samples of the rare class in these subspaces. Interestingly, this approach was inspired in part by work on data augmentation in image data (in the next section [2.3](#)).

The salient difference here is that while image augmentation follows vector operations in most cases, (i.e. shifting, rotating etc.) [SMOTE](#) searches for the nearest neighbors in the feature spaces in the rare class and "interpolates" synthetic samples along the lines between nearest neighbors in the rare class.

Furthermore, in the original paper [Chawla et al. \(2002\)](#) evaluate [SMOTE](#) using another dataset that is of interest, the Pima Indians Diabetes dataset described by [Smith et al. \(1988\)](#) - compare [5](#).

Due to the fact that the objective here is the detection of a rare class, the authors evaluate their approach using the Receiver Operator Characteristic (ROC) Area Under the Curve (AUC) as seen here:

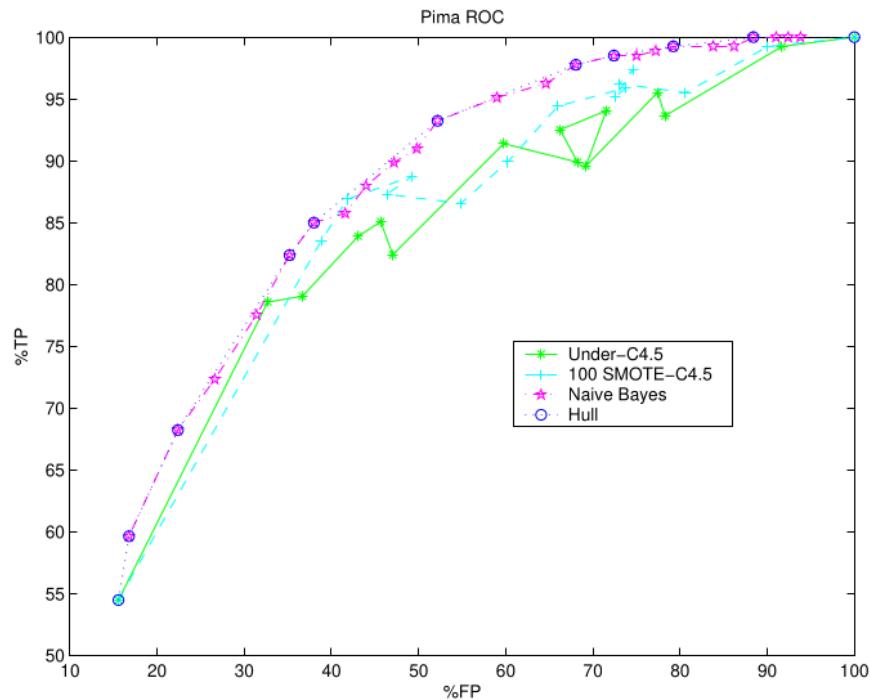


Figure 6: from Chawla et al. (2002)

Comparison to the relevant section here: 5.3.2 Note that while interesting, SMOTE is a strongly opinionated approach that would probably not be suited to application to balanced datasets.

2.3 Data Augmentation for Image Data

While the approaches for synthetic data generation discussed in previous sections involve the addition of noise into the dataset (specifically **SDV** as proposed by [Patki et al. \(2016\)](#) and the interpolation in **SMOTE** proposed by [Chawla et al. \(2002\)](#)), they do not constitute the alteration of existing training data proposed by [Ha and Bunke \(1997\)](#).

On a conceptual level, image data is quite distinct in the way that it includes a notion of locality (neighboring pixels are in relation to one another) and that it lends itself to rapid visual assessment by humans.

In contrast to vectors of attributes (which constitute other datasets), humans are capable of immediately assessing quite literally "at a glance" whether or not the meaning behind an image representation has been preserved after alteration.

This eliminates the need for intermediate tools such as checking for correlation between features etc. (as mentioned specifically in [2.1.2](#)) to assess whether data augmentation techniques yielded qualitative results.

When training **NNs** for image classification therefore, a common practice is **data augmentation**, a range of semi-random transformations applied to images in order to artificially increase the breadth of data that the model is exposed to while presenting the underlying meaning.

Such operations include:

- rotation
- shearing
- zoom
- height & width shift

effectively, these operations transform an image into a completely new observation (from the viewpoint of a [NN](#)) while preserving the underlying signals in the data.

To illustrate the effect, examples can be seen here:

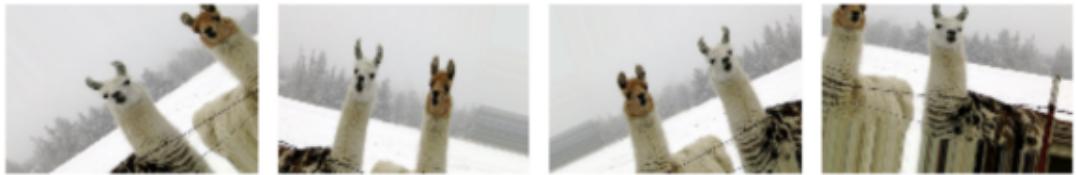


Figure 7: showing classical image transformations from [Perez and Wang \(2017\)](#)

While These transformations seem trivial to a human, it is important to note that effectively all pixels in these images are changed and to a machine learning model, each perturbation represents an entirely new observation. Only the underlying meaning has been preserved, forcing the network to learn hierarchical representations of the subjects in the data, e.g. the presence of ears, eyes furs, their relative positions to one another and so on.

As a note, in this very paper [Perez and Wang \(2017\)](#) mention image augmentation with style-transfer [GANs](#), a currently very active field of study.

However, with other types of data this might not be possible. Attributes of another dataset may not be feasibly 'shifted' in one direction or another without fundamentally changing the signal and misleading the model. Furthermore, humans cannot immediately assess the loss of this meaning in the data, which is why other methods of creating synthetic data (such as [DS](#) in particular [2.1.2](#)) have to rely on statistical indicators of similarity to assess the actual quality of the data.

As alluded to in the previous section, the approach described by [Ha and Bunke \(1997\)](#) partly inspired [SMOTE](#) noted here: [2.2](#).

A large comprehensive survey on the effectiveness on such techniques was composed by [Shorten and Khoshgoftaar \(2019\)](#), comparing the efficacy of a large array of variations.

3. Generative Adversarial Networks

Since the technical implementation chosen here rests on [GANs](#), a brief introduction into the topic will be given here. Currently a very active field of study and experimentation, a history and complete list of applications of [GANs](#) is out of scope of this paper, especially since due to their development on image data, [GANs](#) have been historically quite strongly focused on image data. Some additional thoughts in this regard are found in the section on image data augmentation (here [2.3](#)) and in the discussion (here [6](#)).

Furthermore, many of the points here will be touched upon in the technical implementation (here [4](#)) and will not be repeated here, specifically Wasserstein- and Deep Convolutional GANs ([DCGANs](#)).

3.1 A brief summary of Generative Adversarial Networks

Goodfellow et al. (2014) showed in their original paper proposing the method, that by connecting two NNs; a **generator** and a **discriminator**, these two networks can be induced to engage in a zero-sum competition with each other, thereby teaching the generator to produce realistic synthetic images.

3.1.1 The Generator Network

Receives random input from a latent space, in the form of a vector sampled from a distribution which is either uniform or normally distributed, performs transformations on them and outputs samples with the same dimensionality as the "real" examples.

3.1.2 The Discriminator Network

Receives samples produced by the generator as well as real examples and attempts to predict which are which. Importantly, the inverse of the Discriminator's input constitutes the loss function of the generator, i.e. the generator is learning to "fool" the discriminator.

3.1.3 GAN loss function

The loss function for a GAN as proposed by Goodfellow et al. (2014):

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))]$$

So essentially, what is being optimized for is a composite of:

- the discriminator payoff; $\mathbb{E}_{x \sim p_{data}(x)}[\log D(x)]$; how well the discriminator distinguishes between the samples
- the generator payoff; initially $\mathbb{E}_{z \sim p_z(z)}[\log(D(G(z)))]$; how well the discriminator is fooled by the generator

as noted by Raschka and Mirjalili (2017) the generator payoff is **inverted** in the loss function proposed by Goodfellow et al. (2014).

This is due to the fact that at the very beginning of the training process, the examples produced by the generator are simply noise, and the discriminator is certain which samples are real. Thereby, $\log(D(G(z)))$ (the probability assigned by the discriminator to samples produced by the generator) will be very close to zero, providing a weak signal to the generator (and thus suffering from vanishing gradients).

Therefore in the generator loss function this is inverted to $\mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))]$ and maximized.

Note that this loss function provides diametrically opposed incentives to the two parts of the network, resulting in a **Nash Equilibrium** between the two parts (ideally). GANs are notoriously difficult to balance, which is elaborated on in the technical implementation (here 4.2.2).

3.1.4 GAN training loop

More details on the training loop are mentioned and shown in the relevant section on the technical implementation (here [4.2.2](#)) - briefly however the steps in the training loop are:

- sample random data from the latent space and provide the input to the generator
- create output from the generator, label it automatically, mix it with real data and input into the discriminator
- create discriminator predictions, compute the loss function and adjust the discriminator weights
- provide feedback from the discriminator output to the generator and adjust its' weights

To illustrate, Goodfellow et al. (2014) trained their first GAN on the Deng (2012) dataset on handwritten digits (among others) and showed that the composite network converged on an equilibrium and produced acceptable synthetic handwritten digits:



Figure 8: from Goodfellow et al. (2014)

Further developments are mentioned and implemented in the architecture section of the next chapter (here 4.2.1).

4. Technical Application

4.1 Theoretical applicability

In their landmark paper in 2014, [Goodfellow et al. \(2014\)](#) demonstrated the viability of [GANs](#) on creating image data on the classic MNIST dataset (described by [Deng \(2012\)](#)), by generating - among other things - convincing handwritten digits. As mentioned in [2.3](#), some of the architecture specifics and evaluation are quite specific to image data in that

- the data contains a notion of locality, as neighboring data points (i.e. pixels) are strongly correlated
- dimensionality of the generated data is higher than the **latent space**
- results lend themselves to visual quality inspection by humans (it is easy to see even degrees of quality between different architectures) as explained here ([2.3](#))

specifically the former points are strongly relevant to [GAN](#) architecture, as will become obvious shortly.

4.2 Technical implementation steps

Since the goal of this paper is to evaluate whether or not a **GAN** may be used to not only generate more data of a small non-image dataset (which is fairly trivial) but whether or not this data actually serves to **boost model performance** of models trained on the resulting data, a small, well-understood standard dataset was used to develop the initial architecture; the titanic dataset proposed by [Farag and Hassan \(2018\)](#).

Specifically, the iconic titanic dataset constitutes a binary classification problem, which facilitates quick model evaluation and ameliorates some of the more typical difficulties of training **GANs** - see below.

The first attempts to create a basic, dense **GAN** actually failed to converge for a significant number of experiments with different amounts of layers, neurons and sizes of the latent space.

Somewhat unsurprisingly, achieving the classic Nash Equilibrium between discriminator and generator was fairly difficult and the initial models all proved unstable.

GANs provide several unique challenges, and/or failure modes:

- mode collapse [Che et al. \(2017\)](#)
- oscillation and general instability of the model [Liang et al. \(2018\)](#)
- catastrophic forgetting [McCloskey and Cohen \(1989\)](#)

Mode collapse is especially relevant in a task like MNIST (as used by [Goodfellow et al. \(2014\)](#)), where there are multiple classes to be generated, and the generator becomes increasingly proficient in generating one class explicitly - thankfully, this is less of an issue in a binary classification task.

4.2.1 Network Architecture

The other failure modes, however **did** all make an appearance at one time or another, after the initial data preparation. It was fairly clear that the initial network, with one layer each for the generator and the discriminator each, and 64 neurons had insufficient representational power to converge on creating convincing samples as can be seen in 9:

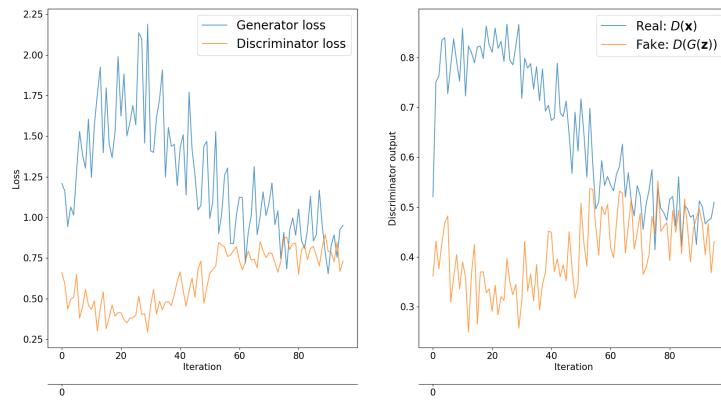


Figure 9: Initial simple dense GAN - left side shows the losses of generator and discriminator, right side shows the probabilities assigned to real and fake samples by the discriminator

Further experiments, with increased numbers of layers and neurons, produced first a very typical oscillation pattern, shown in 10:




Figure 10: Dense GAN, 3 layers, 64 neurons/layer; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator

Note the oscillations in the early stages, which barely decrease in amplitude at all.

Training for an increased number of epochs did not lead to the emergence of a proper equilibrium state, since the network was altogether too unstable. This instability is one of the most typical failure modes in GANs as noted by [Wiatrak et al. \(2020\)](#).

Finally, it has to be stressed that finding the ideal combination learning rates, dropout in the discriminator and number of training epochs, is really quite difficult, especially since there appears no good substitute to visual examination of the pattern that is produced by a given architecture and then adjusting the parameters.

This process that has to be iterated for quite a while, and is fairly manually involved and heavy on trial-and-error.

Ultimately, a promising architecture appeared to be dense networks with 3 layers each, but a higher number of neurons, and still these networks diverged rather quickly shown here 11:




Figure 11: Dense GAN, 3 layers, 128 neurons/layer, reduced learning rate and dropout in discriminator - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator

Unfortunately, from here on out simply optimizing the number of neurons and learning rate and learning rate scheduling was not enough to mitigate this divergence.

Although implementing the popular 1Cycle learning rate decay (described by [Smith \(2018\)](#)) did ameliorate the issue somewhat, it did not fix the network.

What ultimately made the difference is an adaptation of the architecture proposed by Radford et al. (2016). The architecture proposed here for image generation constitutes a **symmetrical** upsampling from the latent space in the generator (in case of images, a **transposed convolution**) and downsampling with the same steps in reverse in the discriminator. As shown by Suh et al. (2019) here:

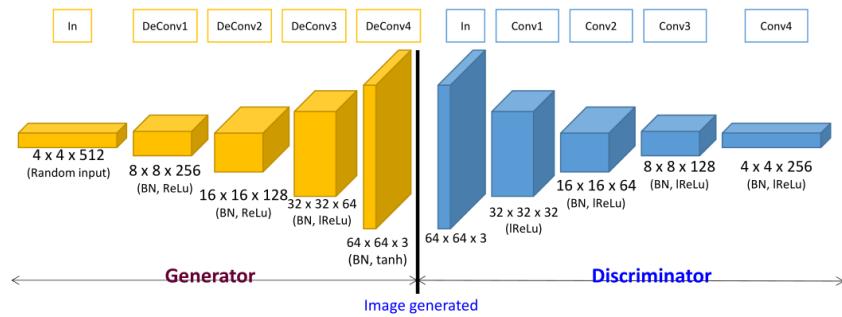


Figure 12: Architecture of Discriminator and Generator

Initially, implementing convolution actually deteriorated performance and completely prevented convergence of the network. One probable explanation could be the fact that convolution and transposed convolution not only downsample and upsample the latent space (respectively) but more fundamentally relate to locality in the data; i.e. multiple convolutional layers over a picture effectively create hierarchical feature extraction - see more details in the section on image data (here 2.3).

A paper that illustrates the mechanics of this fairly well was Dumoulin and Visin (2018). Effectively, these convolutions would initially find small features in images, subsequent convolutions would assemble these features into feature maps and their presence would indicate the presence of objects in an image. The entire concept of strides and adjacent data points however, does not make sense in the concept of a dataset where an observation consists of a feature vector, in which the order of its' attributes does not convey any information.

While 1D convolutions are quite widely used in sequence and time-series processing - which are quite comfortably out of scope of this paper - they fundamentally seem unsuited to a dataset which would not lose any information if the order of its' attributes was permuted.

What **did** make a difference was implementing the symmetry of upsampling and condensing in the generator and discriminator.

Furthermore, Radford et al. (2016) propose other guidelines for building DCGANs which proved helpful:

- implementing BatchNormalization in the generator and discriminator Ioffe and Szegedy (2015)
- using ReLU activation in all layers in the generator except for the output, which would use tanh
- using LeakyReLU in all layers in the discriminator, except for the output which uses sigmoid

After implementing these guidelines, using Binary Categorical Crossentropy loss, the generator and discriminator actually converged fairly well already without diverging quickly, as seen in 13

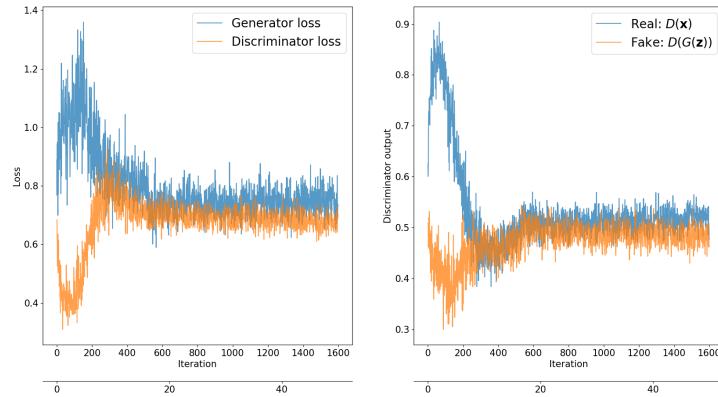


Figure 13: Dense GAN, 2 layers, 32 neurons/layer; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator

Importantly, in this architecture the **generator** model starts with a dense layer containing 32 neurons, which doubles in every layer (once in this case, although this is a variable parameter). The final dense hidden layer is then downsampled again to reflect the original data - like this:

Listing 1: generator network

```

1 import tensorflow as tf
2
3
4 def create_generator_network(
5     number_hidden_layers: int = 2,
6     number_hidden_units_power: int = 5,
7     hidden_activation_function: str = 'ReLU',
8     use_dropout: bool = False,
9     upsampling: bool = True,
10    use_batchnorm: bool = True,
11    dropout_rate: float = 0.3,
12    number_output_units: int = 12,
13    output_activation_function: str = 'tanh') -> tf.keras.Model:
14
15    model = tf.keras.Sequential()
16    for i in range(number_hidden_layers):
17
18        if upsampling:
19            # implements the guideline DCGAN - upsampling layers in the generator
20            model.add(tf.keras.layers.Dense(2 ** (number_hidden_units_power + i), use_bias=False))
21
22    else:

```

```
23         model.add(tf.keras.layers.Dense(2 ** number_hidden_units_power, use_bias=False))
24
25     if use_batchnorm:
26         model.add(tf.keras.layers.BatchNormalization())
27     else:
28         pass
29
30     model.add(tf.keras.layers.Activation(hidden_activation_function))
31
32     if use_dropout:
33         model.add(tf.keras.layers.Dropout(dropout_rate))
34     else:
35         pass
36
37     model.add(tf.keras.layers.Dense(number_output_units))
38     model.add(tf.keras.layers.Activation(output_activation_function))
39
40 return model
```

Listing 1 shows the python function which creates the generator network.

The discriminator implements the exact mirror image of this pattern, beginning with the same amount of neurons after the input layer and downsampling by half each layer;

Listing 2: discriminator network

```

1 import tensorflow as tf
2
3
4 def create_discriminator_network(
5     number_hidden_layers: int = 2,
6     number_hidden_units_power: int = 5,
7     hidden_activation_function: str = 'LeakyReLU',
8     use_dropout: bool = True,
9     upsampling: bool = True,
10    use_batchnorm: bool = True,
11    dropout_rate: float = 0.3,
12    number_output_units: int = 1) -> tf.keras.Model:
13
14     model = tf.keras.Sequential()
15
16     for i in range(number_hidden_layers):
17
18         if upsampling:
19             # implements the guideline - downsample in the discriminator network
20             model.add(tf.keras.layers.Dense(2 ** (number_hidden_units_power + number_hidden_layers - i - 1)))
21
22         else:
23             model.add(tf.keras.layers.Dense(2 ** number_hidden_units_power))
24
25         if use_batchnorm:
26             model.add(tf.keras.layers.BatchNormalization())
27         else:
28             pass
29
30     model.add(tf.keras.layers.Activation(hidden_activation_function))
31
32     if use_dropout:
33         model.add(tf.keras.layers.Dropout(dropout_rate))
34     else:
35         pass
36
37     model.add(tf.keras.layers.Dense(number_output_units, activation=None))
38
39     return model

```

Listing 2 shows the python function which creates the discriminator network.

4.2.2 Network Training Implementation

With the basic architecture for the Network in place, the functions are put together into a custom training loop. While there are multiple approaches to training GANs, starting with the seminal paper by Goodfellow et al. (2014), such as training the discriminator and generator in an epoch separately, here a custom training loop with separate optimizers was chosen from the beginning, in order to accommodate more exotic loss functions (which will become relevant in the section immediately following).

Since a number of experiments on the efficacy of the generated data has to be done, the entire GAN system was set up to be created with sensible defaults, dynamically adapting to 1D datasets. Specifically Buitinck et al. (2013) suggest design lessons from scikit-learn, one of which is using sensible defaults.

In order to automatically create both generators and discriminators dynamically based on input shape (but with strong default settings which ideally do not have to be adjusted during experimentation at all) a small package was created which encapsulates the entire training loop.

Key part is the training loop shown here:

Listing 3: training loop

```

1 # lists to store losses and values
2 all_losses = []
3 all_d_vals = []
4
5 for epoch in range(1, n_epochs+1):
6     epoch_losses, epoch_d_vals = [], []
7     for i,(input_z,input_real) in enumerate(training_data):
8
9         # generator loss, record gradients
10        with tf.GradientTape() as g_tape:
11            g_output = generator_model(input_z)
12            d_logits_fake = discriminator_model(g_output, training=True)
13            labels_real = tf.ones_like(d_logits_fake)
14            g_loss = loss_fn(y_true=labels_real, y_pred=d_logits_fake)
15
16        # get loss derivatives from tape, only for trainable vars, in case of regularization / batchnorm
17        g_grads = g_tape.gradient(g_loss, generator_model.trainable_variables)
18
19        # apply optimizer for generator
20        g_optimizer.apply_gradients(
21            grads_and_vars=zip(g_grads, generator_model.trainable_variables))
22
23        # discriminator loss, gradients
24        with tf.GradientTape() as d_tape:
25            d_logits_real = discriminator_model(input_real, training=True)
26
27            d_labels_real = tf.ones_like(d_logits_real)
28
29            # loss for the real examples - labeled as 1
30            d_loss_real = loss_fn(
31                y_true=d_labels_real, y_pred=d_logits_real)
32
33            # apply discriminator to generator output like a function
34            d_logits_fake = discriminator_model(g_output, training=True)
35
36            # loss for the fakes - labeled as 0
37            d_labels_fake = tf.zeros_like(d_logits_fake)
38
39            # loss function
40            d_loss_fake = loss_fn(
41                y_true=d_labels_fake, y_pred=d_logits_fake)
42
43            # compute component loss for real & fake
44            d_loss = d_loss_real + d_loss_fake
45
46        # get the loss derivatives from the tape
47        d_grads = d_tape.gradient(d_loss, discriminator_model.trainable_variables)
48
49        # apply optimizer to discriminator gradients - only trainable
50        d_optimizer.apply_gradients(
51            grads_and_vars=zip(d_grads, discriminator_model.trainable_variables))
52
53        # add step loss to epoch list
54        epoch_losses.append(
55            (g_loss.numpy(), d_loss.numpy(),
56             d_loss_real.numpy(), d_loss_fake.numpy()))
57
58        # probabilities from logits for predictions, using tf builtin
59        d_probs_real = tf.reduce_mean(tf.sigmoid(d_logits_real))
60        d_probs_fake = tf.reduce_mean(tf.sigmoid(d_logits_fake))
61        epoch_d_vals.append((d_probs_real.numpy(), d_probs_fake.numpy())))

```

```

62     # record loss
63     all_losses.append(epoch_losses)
64     all_d_vals.append(epoch_d_vals)
65     print(
66         'Epoch {:03d} | ET {:.2f} min | Avg Losses >>'
67         ' G/D {:.4f}/{:.4f} [D-Real: {:.4f} D-Fake: {:.4f}]'
68         .format(
69             epoch, (time.time() - start_time)/60,
70             *list(np.mean(all_losses[-1], axis=0))))
71
72     result = {
73         'all_losses': all_losses,
74         'all_d_vals': all_d_vals,
75         'generator': generator_model,
76         'discriminator': discriminator_model}
77
78 if export_generator:
79
80     print()
81     print('saving generator model')
82
83     tf.keras.models.save_model(generator_model, f'./models/generator_{model_name}.h5')
84     print(f'generator model saved to: ./models/{model_name}.h5')
85
86 return result

```

Listing 3 shows the python function which trains the network. Note that this function is **quite strongly simplified**, the actual code used can be found at https://github.com/PaulBFB/master_thesis/blob/main/train_generator.py and would not likely be germane to the paper in its' entirety in any case, even though it proved a strong basis for a large number of experiments in semi-dynamically creating GANs for 1D data.

The above training loop was developed together with the network architecture and also produced the training graphics shown so far.

Before it was used in experimentation however, adjusting its' loss function was tested. Specifically, as proposed by Arjovsky et al. (2017), implementing the Earth Mover's Distance (**EM**) as a loss function.

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

Figure 14: EM / Wasserstein distance formula

As mentioned in the paper, this distance describes the amount of work that is necessary to transform one distribution into another, given an optimal transfer plan γ which actually denotes **how** the work is done. Furthermore, and probably most importantly, the **EM**, in contrast to other loss functions, is actually a function of the parameters θ of the distributions in question, i.e. it can express partial derivatives with respect to the distributions' individual parameters (and thus facilitates Gradient Descent)!

However, finding the ideal transfer plan γ is an optimization problem by itself, since it constitutes an **optimal** solution. As Arjovsky et al. (2017) mention therefore, it is approximated during training.

A complete explication of the metric and its' approximation is out of scope of this paper.

What this achieves in practice, is that it enables the discriminator to act as a **critic**, essentially reporting the distance that the generator has yet to move back during training, which the generator then backpropagates to its' parameters. Thereby, the loss during training actually becomes more meaningfully readable.

[Arjovsky et al. \(2017\)](#) recommend in their paper to clip the gradients reported back to the generator. This led to substantial instability in the network, as can be seen here:

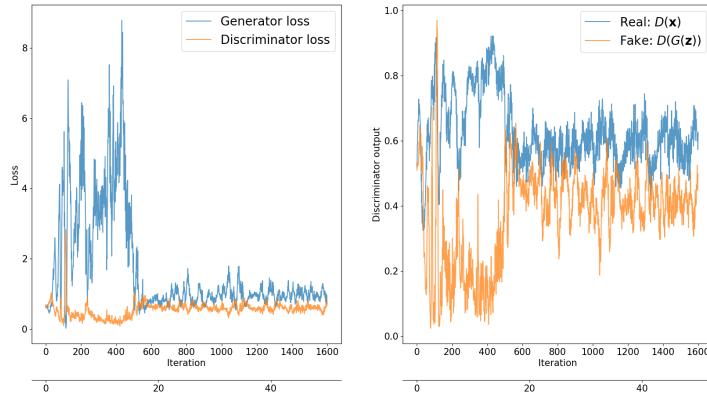


Figure 15: Wasserstein GAN, 3 layers; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator

As can be seen clearly here, the equilibrium between the components is fairly unstable. [Gulrajani et al. \(2017\)](#) pose that gradient clipping here actually leads to exploding and vanishing gradients, which seems to describe this result fairly well.

In the actual implementation therefore, the **Gradient Penalty** method they recommend is implemented, namely:

- between real and fake examples in a batch, choose a random number sampled from a uniform distribution
- interpolate between real and fake examples
- calculate discriminator loss for all interpolated examples
- add the gradient penalty based on the interpolations
- remove batch normalization from the discriminator, since it shifts example gradients based on the entire batch

Therefore, the training loop was modified:

Listing 4: training loop

```

1  for epoch in range(1, n_epochs+1):
2      epoch_losses, epoch_d_vals = [], []
3      for i,(input_z,input_real) in enumerate(training_data):
4
5
6          # set up tapes for both models
7          with tf.GradientTape() as d_tape, tf.GradientTape() as g_tape:
8              g_output = generator_model(input_z, training=True)
9
10         # real and fake part of the critics output
11         d_critics_real = discriminator_model(input_real, training=True)
12         d_critics_fake = discriminator_model(g_output, training=True)
13
14         # generator loss - (reverse of discriminator, to avoid vanishing gradient)
15         g_loss = -tf.math.reduce_mean(d_critics_fake)
16
17         # discriminator losses
18         d_loss_real = -tf.math.reduce_mean(d_critics_real)
19         d_loss_fake = tf.math.reduce_mean(d_critics_fake)
20         d_loss = d_loss_real + d_loss_fake
21
22         # INNER LOOP for gradient penalty based on interpolations
23         with tf.GradientTape() as gp_tape:
24             alpha = rng.uniform(
25                 shape=[d_critics_real.shape[0], 1, 1, 1],
26                 minval=0.0, maxval=1.0)
27
28             # creating the interpolated examples
29             interpolated = (
30                 alpha*tf.cast(input_real, dtype=tf.float32) + (1-alpha)*g_output)
31
32             # force recording of gradients of all interpolations (not created by model)
33             gp_tape.watch(interpolated)
34             d_critics_intp = discriminator_model(interpolated)
35
36             # gradients of the discriminator w. regard to all
37             grads_intp = gp_tape.gradient(
38                 d_critics_intp, [interpolated])[0]
39
40             # regularization
41             grads_intp_l2 = tf.sqrt(
42                 tf.reduce_sum(tf.square(grads_intp), axis=[1, 2, 3]))
43
44             # compute penalty w. lambda hyperparam
45             grad_penalty = tf.reduce_mean(tf.square(grads_intp_l2 - 1.0))
46
47             # add GP to discriminator
48             d_loss = d_loss + lambda_gp*grad_penalty

```

Note that according to [Gulrajani et al. \(2017\)](#) a λ value of 10.0 worked well in all examples, which is what was used here. Again, this is a strongly truncated version of the code, the full version can be found at https://github.com/PaulBFB/master_thesis/blob/main/train_wasserstein_generator.py-this

also contains modified functions to create a **discriminator** without BatchNormalization.

The resulting training loop with the same layers and upsampling-downsampling symmetry between generator and discriminator resulted in this:

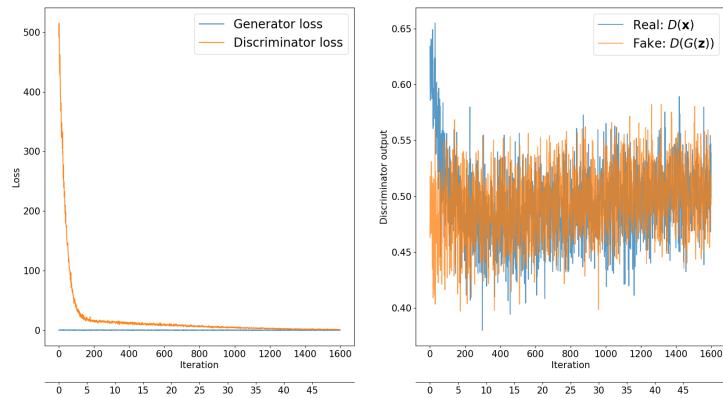


Figure 16: Wasserstein GAN with gradient penalty, 2 layers; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator

The equilibrium is quite tightly clustered around a probability of real and fake samples, which indicates that the discriminator and generator quickly reach a stable equilibrium in which the discriminator is essentially forced to guess between real and fake examples. Note that the probabilities reported by the discriminator cluster very tightly around .5 for both real and fake samples, which overlap almost perfectly.

Both training loops, the basic DCGAN-adapted including BatchNormalization as well as the Wasserstein loop including penalty were preserved as separate modules. In order to perform efficient experiments on entire datasets, again with sensible default values as recommended by [Buitinck et al. \(2013\)](#) those loops were wrapped into a python module, which enables:

- accept preprocessed data in the form of numpy arrays, keep the original size or take a random sample from it (in order to experiment with even smaller data-subsets)
- retrain either a Wasserstein-GP or DCGAN-adapted generator if the data was decreased (or the generator is forced to be trained)
- create a number of samples based on the amount of original data and mix it into the original data

Especially important here is the second point, since using a generator trained on the **full training data set** as a **GAN** and then generating data to mix into a reduced subset would effectively constitute **information leakage** from the entire training dataset!

The complete function enhancing data may be found here https://github.com/PaulBFB/master_thesis/blob/main/enhance_data.py, also the helper function using the generator to generate data from a distribution may be found here: https://github.com/PaulBFB/master_thesis/blob/main/generate_data.py - both functions are fairly basic. Mostly of interest is the fact that they may be used fairly agnostically with a given training dataset, given that it has been processed to be suitable to train standard models on it; i.e. scaled, imputed if necessary and categorically encoded.

From here on out, experiments were performed on the dataset with different types of models.

5. Data Boosting Experiment Results

5.1 Experiments on the original dataset

The technical implementation that has been described in 4 has been entirely developed on the titanic dataset described by Farag and Hassan (2018). Therefore the first experiments were also performed on this dataset.

It is important to **note** here, that due to the fact that this dataset is extremely widely used, there have been significantly higher performances in accuracy achieved. These performances are mostly due to extensive feature engineering, since some of the features of the dataset contain implicit information. Take for example the cabin number, which contains information on where the passenger was staying, which would logically have bearing on their odds of survival, if it was mapped to the cabin's distance from the deck and/or lifeboats.

This is, however, explicitly **not** the purpose of this experiment, a notebook that does this fairly well can be found here: <https://www.kaggle.com/vinothan/titanic-model-with-90-accuracy>

The purpose of this experiment however, is to see if gains in performance can be achieved by simply applying larger compute power to the dataset in an agnostic fashion (more details in the discussion).

5.1.1 Experiments with decreased amounts of data

Initially, the question posed had been whether or not increasing the amount of training data available using a GAN could increase neural network performance. To examine this in detail, the data was systematically decreased in steps, and neural networks were then trained in parallel

- on the shrunken data
- on progressively boosted data

using parameter gridsearch. Gridsearch on neural networks is not yet well automated, therefore in order to do this, some helper modules were created:

A model creation function

To ensure that the models that were trained on the shrunken and boosted data partitions, these models had to be created using the same parameters. This was done with a simple model creation function that encapsulated all the necessary defaults;

Listing 5: training loop

```

1 def make_model(
2     input_shape: tuple = (11, ),
3     number_hidden_layers: int = 8,
4     activation: str = 'elu',
5     alpha: float = .2,
6     neurons: int = 32,
7     loss: str = 'binary_crossentropy',
8     learning_rate: float = .003,
9     dropout_rate: float = .5) -> Model:
10
11     model = models.Sequential()
12     model.add(layers.InputLayer(input_shape=input_shape))
13
14     for i in range(number_hidden_layers):
15         model.add(layers.Dense(
16             neurons,
17             kernel_initializer='he_normal',
18             name=f'hidden_layer_{i}_relu_alpha_{alpha}'))
19
20     if number_hidden_layers >= 3:
21         model.add(layers.BatchNormalization())
22
23     model.add(layers.Activation(activation))
24     model.add(layers.Dropout(dropout_rate, name=f'dropout_{i}_{round(dropout_rate * 100)}'))
25
26     model.add(layers.Dense(1, activation='sigmoid'))
27
28     optimizer = Adam(learning_rate=learning_rate)
29
30     model.compile(loss=loss, optimizer=optimizer, metrics=[accuracy, Precision(), Recall(), AUC()])
31
32     return model

```

The module mostly creates a standard Sequential-class Keras model with minimal dynamic changes (such as BatchNormalization based on the number of layers).

Sklearn-style Gridsearch

Using the KerasClassifier wrapper https://www.tensorflow.org/api_docs/python/tf/keras/wrappers/scikit_learn/KerasClassifier, a basic grid-search function was created to be applied to all models.

Listing 6: gridsearch simplified

```

1  def nn_gridsearch(
2      make_model_function,
3      x_train: np.ndarray = None,
4      y_train: np.ndarray = None,
5      params: dict = None,
6      epochs: int = 100,
7      validation_split: float = .1,
8      patience: int = 10,
9      batch_size: int = 16,
10     n_iterations: int = 10,
11     early_stop: bool = True,
12     save_logs: bool = False,
13     verbose: int = 1):
14
15     keras_cl = KerasClassifier(
16         make_model_function,
17         batch_size=batch_size,
18         shuffle=True,
19         verbose=verbose)
20
21     rnd_search_cv = RandomizedSearchCV(
22         keras_cl,
23         params,
24         n_iter=n_iterations,
25         cv=3,
26         verbose=2,
27         n_jobs=-1)
28
29     callbacks = []
30     if early_stop:
31         callbacks.append(EarlyStopping(patience=patience, monitor='val_loss', mode='min'))
32     if save_logs:
33         callbacks.append(TensorBoard(logdir()))
34
35     rnd_search_cv.fit(
36         x_train, y_train,
37         epochs=epochs,
38         validation_split=validation_split,
39         callbacks=callbacks)
40
41     return rnd_search_cv
42
43
44 if __name__ == '__main__':
45
46     grid_parameters = {
47         'number_hidden_layers': list(range(1, 8)),
48         'neurons': np.arange(1, 100).tolist(),
49         'learning_rate': reciprocal(3e-4, 3e-2).rvs(1000).tolist(),
50         'dropout_rate': np.arange(.2, .6, .1).tolist(),

```

```
51     'alpha': np.arange(.2, .35, .05).tolist(),
52     'activation': ['elu', 'selu', 'relu']}
53
54     grid = nn_gridsearch(
55         make_model,
56         data['x_train_processed'], data['y_train'],
57         grid_parameters,
58         n_iterations=3)
59
60     best_model = grid.best_estimator_.model
61     best_model.save(model_path)
```

Note that both the gridsearch and model function are slightly simplified here, the original is again found at https://github.com/PaulBFB/master_thesis/blob/main/nn_gridsearch.py. Also, to quickly apply the gridsearch function to progressively larger segments of the original data, a script https://github.com/PaulBFB/master_thesis/blob/main/boost_experiments.py was used; which is nothing more than an iteration over the different sized data parts, applying gridsearch to all of them, in order to run in the background (or on remote machines, as some of these tests may take an exceedingly long time, based on the size of the grid and the hardware they run on).

As for the results with decreased data, the following steps were automatically performed:

- shuffle the data randomly, take a small subset from it
- train a generator on it, either a Wasserstein-GP or DCGAN-adapted style generator
- perform gridsearch on it, from the base model creation function
- record the results

All these steps were performed on a subset of the entire data, the training set, the remaining 20% of the data, the test set, was only used for evaluation. Where Gridsearch was applied, a subset of the training data was withheld for Crossvalidation.

The resulting graphic is always formatted in the same way; height of the bar denotes accuracy of the best model, color of the bar denotes the type of boosting that was applied, and the bar position on the x-axis denotes by how much the data was boosted.

Boosting factor ranges are (somewhat arbitrarily chosen, after multiple experiments):

- +20%
- +30%
- +50%
- +200%

For clarity, the accuracy of the unboosted data is marked by a line with its' exact value, to clearly denote performance differences.

Data Size 0.3 - 0.8




Figure 17: Experiment with data sizes 30%, 40% of original data




Figure 18: Experiment with data sizes 50%, 60% of original data

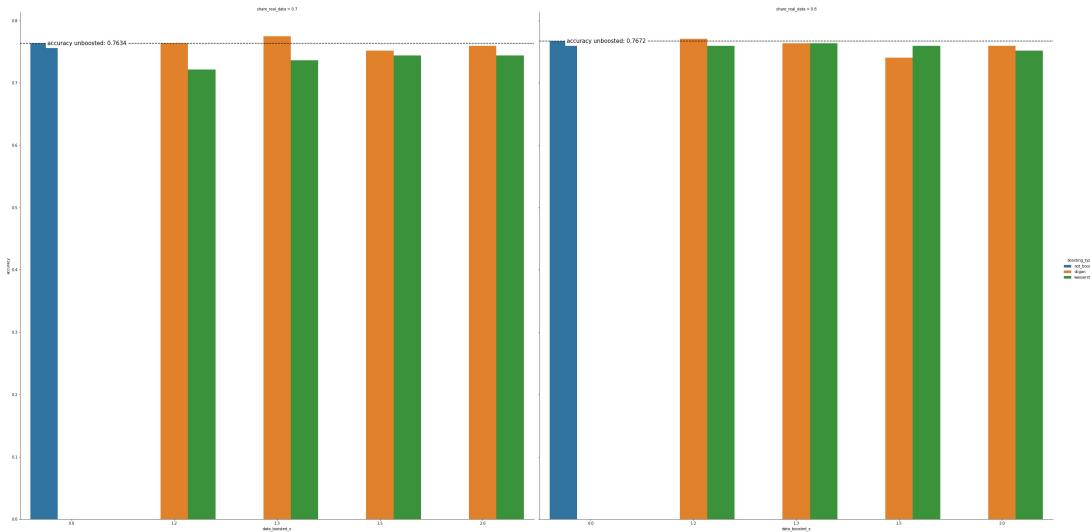


Figure 19: Experiment with data sizes 70%, 80% of original data

As can be seen quite plainly, while at some sizes there appears to be a (very very slight) but static gain in performance, the variation is well within the range of simple random fluctuations due to the random grid search performed on the networks and their random initialization.

Data Size 1

Performing the same test on the entire data:

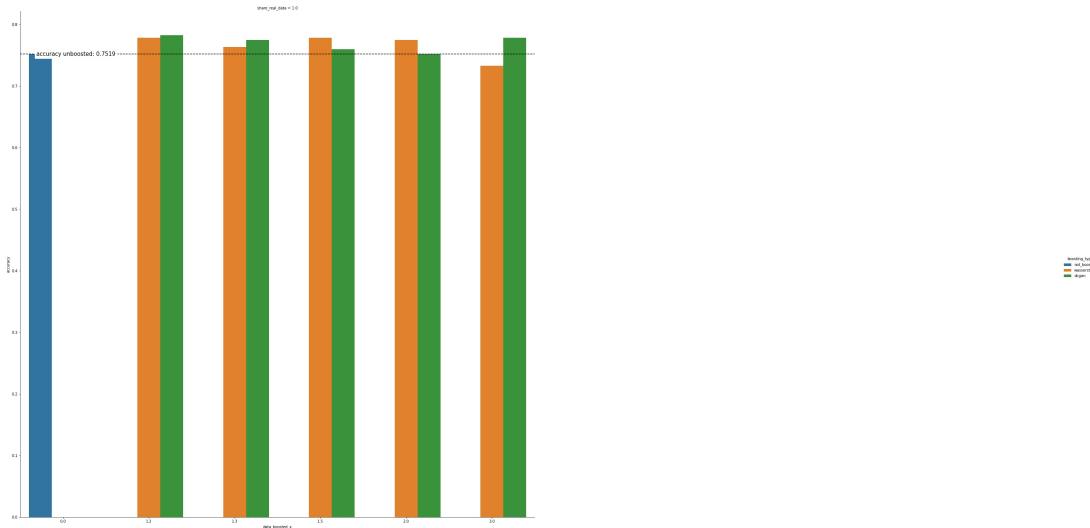


Figure 20: Experiment with data size 100% of original data

appears to yield a performance gain, which is not significant either, but seems more consistent.

Larger Data Sizes

Finally, it was tested whether or not boosting the data to a more significant size would deteriorate performance on these standard models - which would stand to reason.

Any patterns in the training data that do not represent the entire dataset well, would be strongly magnified and therefore skew the result, effectively magnifying the model's generalization error.

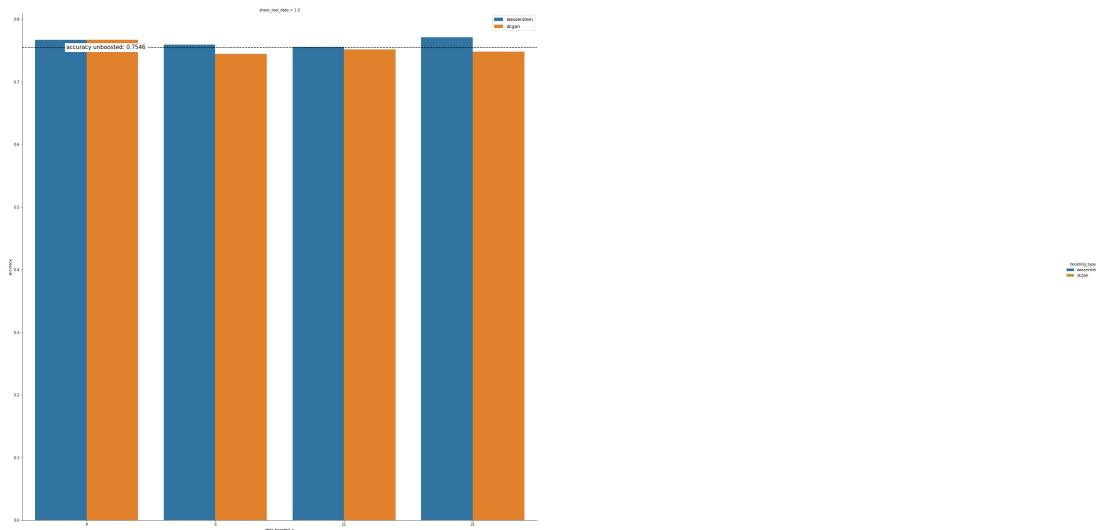


Figure 21: Experiment with data sizes +300% - +2000% original data

The format of the experiment remains the same, the only change being that only models trained on the boosted data are recorded here. The line representing "unboosted" performance takes its' value from the model that was trained on the entire training set.

While the fluctuation in performance remains on the same scale, which quite probably reflects nothing more than random noise, it seems remarkable here that model performance does not, in fact, degrade.

Elaboration on this is found in the final chapter, it just seems prudent to note here that the axiom that comes to mind is:

"if you torture the data long enough, it will confess"

since creating and tuning the entire model architecture and all the helper modules as well as the gridsearch itself was done on the same dataset (so far), to say that this data was tortured is probably quite the understatement.

Therefore, it seemed to be only prudent to use this methodology to test performance on different data sets, as well as its' impact on different types of models.

This accomplished multiple things simultaneously:

- it tested the degree to which fitting the generator and boosting the training data was portable to other small classification datasets
- it tests to what extent the impact on model performance by the boosted data is simply due to an extreme amount of overfitting on the data used so far, as well as information leakage from the training to the test set, due to extensive gridsearch runs.

5.2 Different Model Types and Datasets

In order to compare the impact of this type of boosting on different types of models, a number of standard models were selected from scikit-learn:

- Logistic Regression Classification
- Support Vector Machine Classification
- Random Forest Classification
- Decision Tree Classification
- K-Nearest-Neighbor Classification
- Gaussian Naive Bayes Classification
- Dense Neural Network, for reference

Just a very brief summary of each type of classification (except for neural networks) follows here. All these models were implemented with their **default values** from scikit-learn version 1.0. As noted before, the very sensible default values in scikit-learn are described also by [Buitinck et al. \(2013\)](#) and have served as guidelines of code implemented here. Deep descriptions of these basic models are out of scope of this paper, as they only serve to test the performance of the approach for generation of synthetic training data on new datasets.

The relevant documentation is linked in all sections.

Logistic Regression

Logistic Regression classification, based on scikit-learn, using a liblinear solver as described by [Fan et al. \(2008\)](#).

Logistic Regression calculates a log-probability from the input vector for each observation based on a parameter vector θ . The cost function for misclassification of observations is then normalized by a solver, which constitutes a stochastic gradient descent. Therefore, Logistic Regression in this case behaves similarly to a single layer neural network with one neuron per input parameter.

The regularization that was applied (by scikit-learn default) was ℓ_2 regularization, penalizing parameters.

Documentation can be found here: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html,

Support Vector Machine Classification

Support Vector Machine Classification ([SVC](#)) attempts to formulate a decision boundary around training instances. The goal is to formulate a hyperplane, $n-1$ dimensional space (where n is the number of attributes in the training data) which maximizes the distance to the nearest training instance. Originally formulated by [Vapnik et al. \(1995\)](#) it is a mathematically exceedingly elegant classification solution. In cases where the problem lends itself to projection or dimensionality reduction, [SVCs](#) are known to perform very well.

Documentation found here: <https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html>.

Random Forest & Decision Tree Classification

Decision trees, first described by Quinlan (1986), attempt to iteratively bisect the data in a fashion which increases the resulting halves in terms of their "purity". Purity in this case is described as how "unbalanced" the classes of the target variable are in the resulting halves, the more unbalanced the better (usually tracked in terms of gini impurity or entropy).

By recording these split points, decision trees effectively learn how to halve the data until all resulting observations are in pure subsets (or alternatively a maximum number of splits or a minimum number of samples has been reached).

Decision trees are intuitively well understandable, and are robust against outliers and feature scale, but are prone to overfit training data depending on their depth (the amount of times the classifier is allowed to split the data in training).

Documentation can be found here: <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>

Related to decision trees, **Random Forest Classifiers** train multiple decision trees on different subsets of the data and aggregate their results, and constitute a uniquely adaptable and robust ensemble method.

Documentation can be found here: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>

K-Nearest-Neighbor Classification

K-Nearest-Neighbor Classification ([KNN](#)) calculates all distances between training observations based on their attributes in n-dimensional space (where n is the number of attributes in the training data) and classifies each observation by its' K nearest other observations. The default value of neighbors that is used here is 3 (in order to be less prone to overfitting the training data).

Documentation can be found here: <https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>

Gaussian Naive Bayes Classification

The Naive Bayes Algorithm estimates the probability of a given outcome (1/0 in this case) given the value of a given attribute; as per Bayes' Theorem of conditional probability. Here, the key assumption is that all attributes are normally distributed (hence gaussian) and furthermore that the features are all independent of each other (a strong assumption which makes this model fairly brittle in practice, hence naive).

Treating all attributes independently of one another eliminates the requirement of observing all possible permutations of features for all outcomes.

Documentation can be found here: https://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.GaussianNB.html

5.2.1 Reference - model fitting on the titanic dataset

All these models have been fit on the titanic dataset for reference. In order to make these models comparable, they have all been fit on all datasets with the same parameters (as mentioned in the sections) as well as using 80% of their data as training data, withholding 20% for performance testing.

All training data have then been used to fit a Wasserstein-GP generator and increase their training data size by 20% (since this seemed to be the most consistent and promising configuration).

For clarity, the baseline for a "naive" prediction (that is, always simply predicting the most common class) is also shown, in order to demonstrate whether any of the models exhibit predictive power beyond simply making the safest guess (i.e. a static model).




Figure 22: Comparing all model type performances, training data size boosted by 20% with Wasserstein-GP generator

The outcome here seems quite consistent, with no change in most models, deterioration of performance does not actually occur, which is promising. Given the fact that this dataset is the one the Wasserstein-GP architecture was built on, any further deductions seem spurious.

5.3 Fitting the Generator and models on completely new data

Two new datasets were selected as test sets for the architecture that was used:

5.3.1 The Wine Quality Dataset

Described by [Cortez et al. \(2009\)](#). The dataset contains 6000+ observations of wine quality with 11 attributes and one target variable (quality). The attributes describe characteristics such as citric acid, chlorides, residual sugar, alcohol...

Data preparation steps (using sklearn Pipelines):

- target variable (quality) changed to binary (any quality over 6 is considered quality wine)
- imputation of median value into columns with missing data ("fixed acidity", "pH", "volatile acidity", "sulphates", "citric acid", "residual sugar", "chlorides", number of missing values <10)
- OneHotEncoding of categorical attribute "wine type"
- MinMaxScaling of numerical attributes

Again, feature engineering of the data was deliberately excluded, just as with the titanic dataset. The steps taken are the bare minimum in order to enable machine learning at all by design. As a final note, the dataset used here is fairly unbalanced, roughly 80% falling into the more common class (low quality wine). More on this in the discussion.

Fitting the Wasserstein-GP model on this data, after preprocessing but without changing any of the default values (layers, number of epochs, architecture, loss

function, dynamically building the models from the input shape) yielded the following training log:




Figure 23: Wasserstein GAN with gradient penalty, 2 layers; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator

Encouragingly, the loss for both modules seems to decrease in much the same way as on the original dataset, and the equilibrium denoted by real and generated data clustering together, each at around 0.5 probability by the discriminator seems to hold fairly stable as well, even though the dataset has different size and shape from the original.

Note: details on the training module found here: [4.2.2](#)

The result of all boosted models, again using default values, analogous to the comparison on the titanic dataset:




Figure 24: Comparing all model type performances, training data size boosted by 20% with Wasserstein-GP generator

After the generator was trained on the entire training data (withholding a test set completely to avoid information leakage), in the interest of consistency with the original model, the training data was boosted by 20%.

Now immediately apparent here seems the fact that most models barely exhibit any predictive power "out of the box" at all - which seems fair, given that there appears to be a rare class problem.

The [NN](#) seems to be fairly robust against this, even gaining some predictive power from boosting the data. Also, the models with the biggest (negative)

impact in performance appear to be the most brittle models, Decision Tree and Naive Bayes.

Promisingly, the other models (while not gaining performance) appear to barely be affected at all, which seems to indicate that the generated data overall reflects the pattern in the training data.

A more detailed analysis follows in the discussion.

5.3.2 The Pima Indians Diabetes Dataset

A dataset originally created by the National Institute of Diabetes and Digestive and Kidney Diseases in order to classify diabetes in female patients.

The dataset includes 8 diagnostic markers:

- Pregnancies
- Glucose
- BloodPressure
- SkinThickness
- Insulin
- BMI
- DiabetesPedigreeFunction
- Age

As well as an outcome column.

Initially described by [Smith et al. \(1988\)](#), classifying diabetes using neural networks.

The dataset is fairly balanced, with around 65% of outcomes in the more common class, making it fairly similar to the titanic dataset in this regard (63%, for comparison).

Data preparation steps (using sklearn Pipelines):

- Imputation of missing values
- smoothing of extreme outliers (.05 - .95 percentiles)
- MinMaxScaling of numerical attributes

Again, extensive preprocessing or feature engineering was excluded. The steps taken are the bare minimum in order to enable machine learning at all by design.

The Wasserstein-GP model was fit automatically on the data, after preprocessing but without changing any of the default values (layers, number of epochs, architecture, loss function, dynamically building the models from the input shape) yielded the following training log:




Figure 25: Wasserstein GAN with gradient penalty, 2 layers; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator

Once again, even with the data being quite similar to the original dataset the architecture was built on, the components seem to establish an acceptable equilibrium quickly, which also appears to be fairly stable. With more details in the discussion, this seems to point towards the architecture being fundamentally sensible.

After training the generator agnostically i.e. out of the box only using the default values, the training data was boosted by 20% and the standard models were trained on the data for comparison:




Figure 26: Comparing all model type performances, training data size boosted by 20% with Wasserstein-GP generator

The impact of inserting synthetic data seems to be less pronounced than on the performance of classifiers in the wine quality dataset (24) which may be due to the structural similarities between this dataset and the titanic dataset (specifically its' class balance). A detailed discussion in the final chapter.

5.4 Replacing Training Data with Synthetic Data

Finally, since it seems pertinent, it was tested what the effect of **entirely replacing the training data** would have on the performance of models trained on entirely synthetic data.

To test this, the following steps were performed on the titanic, diabetes and wine datasets:

- splitting the data into training and test data (20% withheld)
- fitting a Wasserstein-GP [GAN](#) on the training data
- completely discard the real training data and replacing it with synthetic data of identical size

Note that both the attributes of the data as well as the target attribute was replaced with synthetic data.

5.4.1 Models on purely synthetic Titanic Data




Figure 27: Comparing all model type performances, training data entirely replaced by synthetic data of same size with Wasserstein-GP generator

Obviously, the effect on model performance is fairly stark. As a note, the line represents a static model, always predicting the most common class in the training data. Given, some models seem to preserve a small amount of predictive power above purely guessing, but this is negligible.

5.4.2 Models on purely synthetic Wine Data

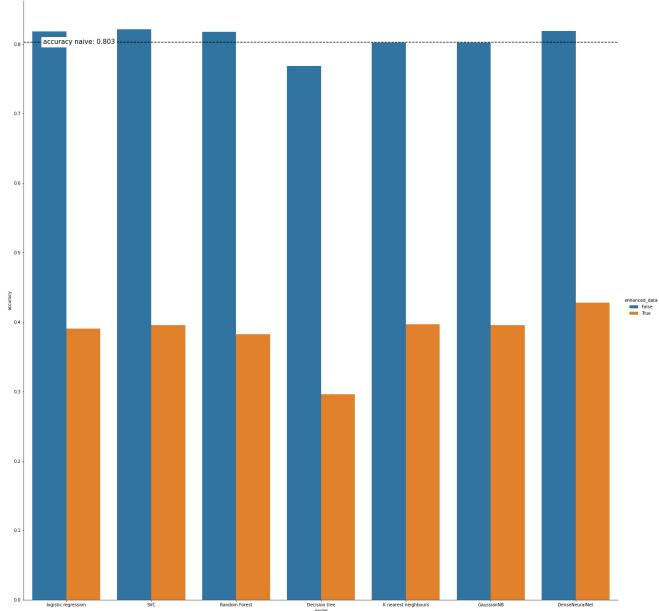


Figure 28: Comparing all model type performances, training data entirely replaced by synthetic data of same size with Wasserstein-GP generator

Quite similarly to the titanic dataset, the effect on predictive power is devastating. Given that the models all barely outperformed a naive majority classifier this is fairly disappointing.

Interesting to note here seems to be the marked symmetry between model classes trained on the original data and the replaced data. Also, the fact that all models created would actually perform better if they were to be inverted is almost an achievement; dubiously epitomized by the decision tree classifier.

Maybe the fact that the training log of the Wasserstein-GP model (shown here [23](#)) appears to show a slight drift out of equilibrium in the later epochs is a contributing factor to this effect.

5.4.3 Models on purely synthetic Diabetes Data




Figure 29: Comparing all model type performances, training data entirely replaced by synthetic data of same size with Wasserstein-GP generator

The analogous experiments were performed with the diabetes dataset, and while the performance decreases markedly, it does so substantially less so than in relation to the wine dataset.

It stands to reason that this hints at a fundamentally better fit of the Wasserstein-GP architecture chosen to the diabetes dataset, probably due to its' structural similarity to the titanic dataset.

5.4.4 Replacement Conclusions

Two interesting notes here, firstly the model class that preserves the most predictive power is the Random Forest Classifier, which given the fact that it constitutes an ensemble model is not overly surprising. Secondly, the stark drop in performance of the NN is - in part - due to the fact how it is trained;

In lieu of more traditional regularization techniques, the model uses keras's EarlyStopping Callback https://keras.io/api/callbacks/early_stopping/ which monitors a key metric and stops the training if the model does not improve for a set number of epochs. In the case of the model used here, the metric used for stopping was the validation loss, calculated on a subset of the training data during the epoch (20%, randomly chosen).

The code for the sample neural network:

Listing 7: base model

```

1 from tensorflow.keras import models, layers
2 from tensorflow.keras.callbacks import EarlyStopping, TensorBoard
3 from nn_gridsearch import logdir
4
5
6 model = models.Sequential()
7 model.add(layers.InputLayer(input_shape=x_train_processed.shape[1:]))
8
9 for i in range(5):
10     model.add(layers.Dense(
11         64,
12         kernel_initializer='he_normal',
13         name=f'hidden_layer_{i}'))
14     model.add(layers.BatchNormalization())
15     model.add(layers.Activation('selu'))
16     model.add(layers.Dropout(0.3, name=f'dropout_{i}_30'))
17
18 model.add(layers.Dense(1, activation='sigmoid'))
19
20 #model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])
21 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
22
23 history = model.fit(
24     enh['x_train_processed'], enh['y_train'], validation_split=.2, epochs=300,
25     callbacks=[

26         EarlyStopping(patience=20, monitor='val_accuracy', mode='max', restore_best_weights=True),
27         TensorBoard(logdir(hyperparam_note='titanic_replacement'))])

```

Listing 7 shows the python function which creates the base neural network

used in all experiments.

While this is beneficial in a standard sequential model, if the trained data is not closely enough representing all data, this drift will be greatly magnified.

Taking a look at the training log (taken from TensorBoard) seems to confirm this:

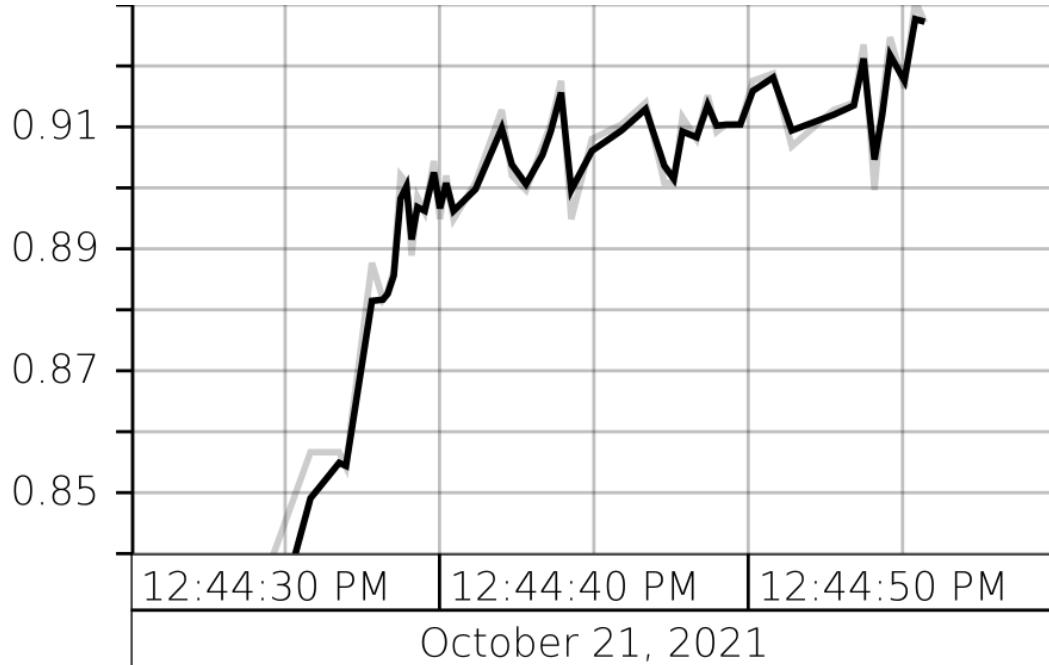


Figure 30: model accuracy history during training; data entirely replaced by synthetic data of same size with Wasserstein-GP generator

So it would seem from this training log, that the accuracy recorded during training quickly increases above 90%. As a side note, with the titanic dataset the generalization error, that is the difference between validation accuracy calculated on the withheld subset of training data in each epoch and accuracy on the test set was always around 3-4 percentage points.

Ultimately, if the synthetic data **were** to be representative of the test data, an accuracy of this magnitude with the same architecture would not be possible. There is a strong chance then, that the network trained here is actually learning some of the patterns that the generator uses to create new data from the random distribution in its' latent space.

6. Discussion

6.1 General Thoughts

Here the general outcomes from every aspect of the thesis will be discussed.

The initial motivation was to evaluate whether data augmentation of small datasets with [GANs](#) is a feasible low-barrier method to increasing model performance, maybe even an alternative to collecting additional data and/or extensive feature engineering.

While algorithmic approaches such as [bagging](#), as described by [Breiman \(1994\)](#) offer substantial performance boosts, it seems that even an apparently stable [GAN](#) - while able to produce superficially similar examples to the initial data distribution - does not actually help to extract additional [meaning](#) from small sets of 1-dimensional data.

There's a quote, probably apocryphal but fairly apropos here;

"I have not failed. I've just found 10,000 ways that won't work."

While certainly not thousands of ways, quite a few were found.

However, as will be noted in the following sections, some of the results may still be of value. Specifically I suspect the adaptation of the symmetrical [GAN](#)

upsampling-downsampling architecture of generator and discriminator may find some use.

6.2 Conclusions from architecture and technical design

First of all, it really cannot be stressed enough how helpful and impactful the design principles in scikit-learn described by [Buitinck et al. \(2013\)](#) were.

Both in the implementation of the experiments in repeatedly shrinking and boosting the data, as well as in the module implementing neural network gridsearch "bridging" scikit-learn and keras (here [6](#)) using these guidelines substantially accelerated the development and experimentation process since it made the practice of repeatedly creating the standard [GAN](#), boosting data and mixing it into existing training data much quicker.

The guidelines also served as a sort of guardrail to design experiment processes along of.

Successfully creating a [GAN](#) system without any additional parameters that appears to - somewhat - establish equilibrium on arbitrary new training data was certainly a helpful milestone.

GAN architecture itself was comfortably the biggest challenge during the project, simply because

- one can never be certain whether or not the network is not converging because the correct configuration of
 - layers
 - neurons
 - learning rate / learning rate scheduling
 - activation function(s)
 - dimension and type (uniform versus normal) of the latent space
 - number of epochs

has simply not been found yet, or if there is no such architecture for just this case

- whether or not the network has established a semi-stable equilibrium yet is difficult to check for automatically (at least, no universal method has been established in the literature)
- the quality of non-image data generated by a **GAN** is generally difficult to assess beyond a simple check whether or not the values fall within the parameters of the attribute that is being modeled.

To expand on the last point, since a lot of literature concerning Generative Adversarial Networks is focused on image generation, a task in which humans still have largely unchallenged **general** capacity, a common identification of failure modes (such as mode collapse described by [Che et al. \(2017\)](#)) is simply "look at it".

In the case of data in the form of attribute vectors it becomes fairly intractable to quickly and informally assess whether the generator network has been stuck in generating a certain subset of the sampling space.

This is actually a critical point; in the example of data generation on new datasets, it quickly became obvious that balance of the dataset is a key factor in model performance. This intuitively makes sense, since a generator that was stuck in a subset of the vector space would wildly unbalance any training data that is synthetically increased with it or increase an existing slight imbalance.

There may be applications in using Conditional **GAN**s as described by [Mirza and Osindero \(2014\)](#) in order to create stratified synthetic training data; an approach that would actually be somewhat analogous to **SMOTE** as described by [Chawla et al. \(2002\)](#).

Using multiple wrapper modules to generate different data shrink and boost sizes, run gridsearch on them and record the results (as well as any resulting **GAN** generator training logs) proved as an involved but ultimately fairly efficient process to thoroughly evaluate the effect on predictive power.

6.2.1 Gridsearch

Specifically concerning gridsearch performed on neural networks using the scikit-learn wrapper as mentioned here [6](#), while generally a promising design pattern that can definitely be reused, ultimately proved to be fairly time consuming (due to the duration of the experiments, fitting a large number of networks using Crossvalidation and selecting the best one).

Furthermore, a byproduct of the specific gridsearch chosen, randomized grid-search, was some amount of randomness in the ultimate performance of the models, which made interpretation of these results more difficult.

This is also the reason it was not actually applied to the new datasets, as it would likely do nothing but obfuscate the results - which are now thankfully (if personally disappointingly) plain.

6.3 Conclusions from shrinking and boosting the original dataset

In this section, the approach that was chosen actually showed the most promise.

Interestingly enough, model performance never degraded at all. While it was strange to see that the model hardly suffered from decreased data size at all, the results on all datasets here seem to indicate that this [GAN](#) architecture is aligned to the specific dataset to such a degree as to not degrade model performance, when it is used to augment either subsets of the data or the entire training data - in which case it may be even beneficial.

A **strong** caveat here is, again, the fact that this data was used to create the architecture so information leakage can hardly be avoided.

It is also possible that models trained on a boosted dataset here learned to ignore the boosted data by identifying patterns created by the generator.

This is only speculation, but an indicator seems to be the fact that in experiments with significantly larger data boosts (up to +20x, see [21](#)) model performance did not decrease, but with complete replacement (see [27](#)) model performance did in fact decrease significantly, if less so than with other datasets.

Although as noted in the section on replacement conclusions, part of the effect, specifically on [NNs](#) is most likely due to the nature of the training process

A passable conjecture here is a composite effect of multiple individual effects, that is:

- performance deterioration due to the training specifics used in the [NN](#) as seen here [30](#)
- differences in representation power between real and synthetic training data
- the fact that gridsearch, by searching a large hyperparameter space, is uniquely likely to result in information leakage (and ignoring synthetic training data)

The fact remains, that on the original dataset, with a large amount of manual [GAN](#) design and searching for ideal hyperparameters, in the right circumstances **slight improvements** may be possible.

6.4 Conclusions from application to other datasets

The entire approach of selecting a mini-battery of models using default parameters and applying them to datasets without any feature engineering with and without boosted data seemed essential after the amount of time, architecture optimization (of the generator) and processing power (on finding the ideal model architecture) that had been applied to the original dataset.

After all of this, there was simply no way that information leakage (and by extension, overfitting on the test data) had not occurred, **especially** the dataset in question was small by design and the model applied to it had a fairly outsized amount of representational power.

While this was fairly intentional since the explicit goal was to evaluate whether using more complex models (i.e. [NNs](#)) on such datasets by increasing their size was a feasible trade-off against feature engineering. This would have required a gain in performance, however.

Thankfully, a large and diverse amount of machine learning datasets are now publicly available, and the additional work of selecting new datasets as a "blank canvas" to perform a sanity check on is tractable.

As a side note, before the explosion of machine learning popularity and the resulting immense increase in availability of resources online this would not have been possible at all, simply due to the fact that sufficiently analogous datasets were not to be found.

In terms of applying the modules that were developed to completely new data (without any configuration or adaptation of the architecture of the [GAN](#)) the result could not have been more promising.

The generator logs that are included (for wine classification here [23](#) and diabetes here [25](#)) were genuinely produced with the finished data enhancement module completely out of the box.

Especially after the initial difficulties of creating any architecture that would converge on vector data at all, and the resulting effort involved in creating the training loop and adding functionality for Wasserstein-GP style generator training, to say this was promising would be an understatement.

The training logs for new data did show some instability though, and from the detailed results of the standard models applied to unchanged and boosted data of these datasets, some have become quite clear:

- the degree to which the dataset is **analogous** to the original dataset makes a significant difference. Specifically the fact that the wine quality dataset has a class balance of roughly 1-4 in contrast to the titanic and diabetes datasets with roughly 1-2 seemed to make a large impact
- any data from an automatic [GAN](#) architecture that is not well tuned to the specific training data will immediately degrade performance

Concerning the second part, I **suspect** that this has to do specifically with the ratio of the size of the latent space to the attributes of the training data and the starting number of neurons (before the first upsampling in the generator).

This is purely speculation at this point, but for the initial architecture convergence was completely impossible until the latent space was set at a standard dimension of (8,). The titanic dataset has 9 attributes and the initial layer of

the generator has 32 neurons by default. The fact that there is a good degree of symmetry here may be a contributing factor.

This hypothesis would be difficult to confirm though, without finding a large number of similar and dissimilar datasets and then evaluating the approach on them. Even then, the effect might be too small to completely confirm or deny it with confidence.

6.5 Conclusions from complete data replacement

As discussed, there are plenty of more proven methods of synthetic data generation, some of which offer model performance which is fairly consistent with the original data already, as discussed here [Hittmeir et al. \(2019\)](#).

While using this method of data generation to create anonymous data was decidedly not the goal from the start, training models on completely synthetic data as mentioned in the paper offers a useful benchmark whether or not the generated data is actually meaningfully reflective of the original data at all.

If we consider the extraction of a model out of a training data set as an efficient representation of the information contained in the data with regard to the target variable, the degradation of the predictive power of the same model with respect to the target variable on synthetic data can be constructed as a useful indicator whether or not the actual meaning of the data was preserved.

Also, and even more acutely relevant, after performing gridsearch on the original data and training models on synthetically increased datasets other than titanic, it was still unclear whether the models trained on such synthetic data retained some predictive power simply due to the degree to which they managed to identify (and ignore) the synthetic data and thereby effectively acting in part like a pseudo-discriminator.

Concerning the specific results of completely synthetic data model performance, I suspect that this is likely - especially considering the models that preserved the most predictive power with partly synthetic data were random forest models, which are notoriously robust (due to them being ensemble models).

Perhaps most interestingly here, all things considered, is the fact that the diabetes dataset experiences the smallest relative loss of predictive power (here:

[29](#)), in relative terms, when data has been completely replaced by synthetic data. Naively, it would make more sense to see the dataset the architecture was originally fit on to experience the smallest loss (actual loss seen here: [27](#)).

While the fact that data generation with a such a [GAN](#) is an intriguing option, especially since it does not rely on any assessment of similarity to generated data in an iterative fashion as the [SDV](#) described by [Patki et al. \(2016\)](#) or the [DS](#) described by [Ping et al. \(2017\)](#) and is a much more unsupervised method, its' results are not in any way comparable. In the future, there may be a way to use conditional [GANs](#) as described by [Mirza and Osindero \(2014\)](#) to achieve the best of both worlds.

6.6 Final Thoughts

Ultimately, some parts of the approach taken show promise and may even be reused, specifically:

- the [GAN](#) architecture pattern using up-downsampling
- the Wasserstein-GP training loop
- the [NN](#) gridsearch

However, the key issue seems to be the fit between the exact [GAN](#) architecture and the specific dataset.

The fact that [GAN](#) training is not only notoriously difficult but also difficult to automate and the quality of vector data generated by them may not be feasible to be automatically assessed, completely automating this approach using the guidelines proposed by [Buitinck et al. \(2013\)](#) in scikit-learn is probably not feasible.

The purpose of this paper was mainly to show the trade-off between data scientist time investment (in feature engineering, data enrichment and - worst case - even data labeling) and increased processing power; i.e. using outsized representational power and computation to draw increased meaning from an otherwise unchanged dataset, as it is done in approaches like [bagging](#).

What appears to be the case here is that the time investment using this approach would - in the ideal case - mainly shift from feature engineering to [GAN](#) architecture search.

In cases where only a small amount of data is available and increasing the training dataset is otherwise impossible **and** feature engineering and data enrichment has been completely exhausted, the approach may hold some promise.

Whether or not the ideal architecture may be found dynamically (as mentioned here 6.4) remains an open question.

As for now, there appears to be no free lunch available.

Bibliography

- Anguita, D., Ghio, A., Ridella, S., and Sterpi, D. (2009). K-fold cross validation for error rate estimate in support vector machines. In *DMIN*, pages 291–297.
- Arjovsky, M., Chintala, S., and Bottou, L. (2017). Wasserstein gan.
- Breiman, L. (1994). Bagging predictors.
- Brownlee, J. (2020). *Data preparation for machine learning: data cleaning, feature selection, and data transforms in Python*. Machine Learning Mastery.
- Buitinck, L., Louppe, G., Blondel, M., Pedregosa, F., Mueller, A., Grisel, O., Niculae, V., Prettenhofer, P., Gramfort, A., Grobler, J., Layton, R., Vanderplas, J., Joly, A., Holt, B., and Varoquaux, G. (2013). Api design for machine learning software: experiences from the scikit-learn project.
- Chawla, N. V., Bowyer, K. W., Hall, L. O., and Kegelmeyer, W. P. (2002). Smote: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16:321–357.
- Che, T., Li, Y., Jacob, A. P., Bengio, Y., and Li, W. (2017). Mode regularized generative adversarial networks.
- Chollet, F. (2017). *Deep learning with Python*. Simon and Schuster.
- Cortez, P., Cerdeira, A., Almeida, F., Matos, T., and Reis, J. (2009). Modeling wine preferences by data mining from physicochemical properties. *Decision*

- Support Systems*, 47(4):547–553. Smart Business Networks: Concepts and Empirical Evidence.
- Deng, L. (2012). The mnist database of handwritten digit images for machine learning research. *IEEE Signal Processing Magazine*, 29(6):141–142.
- Dietterich, T. G. (2000). Ensemble methods in machine learning. In *International workshop on multiple classifier systems*, pages 1–15. Springer.
- Dumoulin, V. and Visin, F. (2018). A guide to convolution arithmetic for deep learning.
- El Emam, K., Mosquera, L., and Hoptroff, R. (2020). *Practical Synthetic Data Generation: Balancing Privacy and the Broad Availability of Data*. O'Reilly Media.
- Fan, R.-E., Chang, K.-W., Hsieh, C.-J., Wang, X.-R., and Lin, C.-J. (2008). Liblinear: A library for large linear classification. *the Journal of machine Learning research*, 9:1871–1874.
- Farag, N. and Hassan, G. (2018). Predicting the survivors of the titanic kaggle, machine learning from disaster. In *Proceedings of the 7th International Conference on Software and Information Engineering, ICSIE '18*, page 32–37, New York, NY, USA. Association for Computing Machinery.
- Géron, A. (2019). *Hands-on machine learning with Scikit-Learn, Keras, and TensorFlow: Concepts, tools, and techniques to build intelligent systems*. O'Reilly Media.
- Goodfellow, I. J., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., and Bengio, Y. (2014). Generative adversarial networks.
- Gulrajani, I., Ahmed, F., Arjovsky, M., Dumoulin, V., and Courville, A. (2017). Improved training of wasserstein gans.
- Ha, T. and Bunke, H. (1997). Off-line, handwritten numeral recognition by perturbation method. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 19(5):535–539.

- Hittmeir, M., Ekelhart, A., and Mayer, R. (2019). On the utility of synthetic data: An empirical evaluation on machine learning tasks. In *Proceedings of the 14th International Conference on Availability, Reliability and Security*, ARES '19, New York, NY, USA. Association for Computing Machinery.
- Ioffe, S. and Szegedy, C. (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift.
- Liang, K. J., Li, C., Wang, G., and Carin, L. (2018). Generative adversarial network training is a continual learning problem.
- McCloskey, M. and Cohen, N. J. (1989). Catastrophic interference in connectionist networks: The sequential learning problem. In *Psychology of learning and motivation*, volume 24, pages 109–165. Elsevier.
- Mirza, M. and Osindero, S. (2014). Conditional generative adversarial nets. *CoRR*, abs/1411.1784.
- Munson, M. A. (2012). A study on the importance of and time spent on different modeling steps. *SIGKDD Explor. Newsl.*, 13(2):65–71.
- Patki, N., Wedge, R., and Veeramachaneni, K. (2016). The synthetic data vault. In *2016 IEEE International Conference on Data Science and Advanced Analytics (DSAA)*, pages 399–410. IEEE.
- Pereira, F., Norvig, P., and Halevy, A. (2009). The unreasonable effectiveness of data. *IEEE Intelligent Systems*, 24(02):8–12.
- Perez, L. and Wang, J. (2017). The effectiveness of data augmentation in image classification using deep learning. *arXiv preprint arXiv:1712.04621*.
- Ping, H., Stoyanovich, J., and Howe, B. (2017). Datasynthesizer: Privacy-preserving synthetic datasets. In *Proceedings of the 29th International Conference on Scientific and Statistical Database Management*, pages 1–5.
- Quinlan, J. R. (1986). Induction of decision trees. *Machine learning*, 1(1):81–106.

- Radford, A., Metz, L., and Chintala, S. (2016). Unsupervised representation learning with deep convolutional generative adversarial networks.
- Raschka, S. and Mirjalili, V. (2017). Python machine learning: Machine learning and deep learning with python. *Scikit-Learn, and TensorFlow. Second edition ed.*
- Raudys, S. J., Jain, A. K., et al. (1991). Small sample size effects in statistical pattern recognition: Recommendations for practitioners. *IEEE Transactions on pattern analysis and machine intelligence*, 13(3):252–264.
- Shorten, C. and Khoshgoftaar, T. M. (2019). A survey on image data augmentation for deep learning. *Journal of Big Data*, 6(1):1–48.
- Smith, J. W., Everhart, J. E., Dickson, W., Knowler, W. C., and Johannes, R. S. (1988). Using the adap learning algorithm to forecast the onset of diabetes mellitus. In *Proceedings of the annual symposium on computer application in medical care*, page 261. American Medical Informatics Association.
- Smith, L. N. (2018). A disciplined approach to neural network hyperparameters: Part 1 – learning rate, batch size, momentum, and weight decay.
- Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., and Salakhutdinov, R. (2014). Dropout: a simple way to prevent neural networks from overfitting. *The journal of machine learning research*, 15(1):1929–1958.
- Suh, S., Lee, H., Jo, J., Lukowicz, P., and Lee, Y. (2019). Generative oversampling method for imbalanced data on bearing fault detection and diagnosis. *Applied Sciences*, 9:746.
- Swingler, K. (1996). *Applying neural networks: a practical guide*. Morgan Kaufmann.
- Vapnik, V., Guyon, I., and Hastie, T. (1995). Support vector machines. *Mach. Learn*, 20(3):273–297.

Wiatrak, M., Albrecht, S. V., and Nystrom, A. (2020). Stabilizing generative adversarial networks: A survey.

Wigner, E. P. (1990). The unreasonable effectiveness of mathematics in the natural sciences. In *Mathematics and Science*, pages 291–306. World Scientific.

List of Acronyms

CNN Convolutional Neural Network

GB Gradient Boosting

NN Neural Network

ML Machine Learning

SMOTE synthetic minority oversampling technique

GAN Generative Adversarial Network

DCGAN Deep Convolutional GAN

EM Earth Mover's Distance

SVC Support Vector Machine Classification

KNN K-Nearest-Neighbor Classification

bagging bootstrap aggregating

GDPR General Data Protection Regulation

HIPAA US Health Insurance Portability and Accountability Act

API Application Programming Interface

SDV Synthetic Data Vault

DS DataSynthesizer

SVM Support Vector Machine

ROC Receiver Operator Characteristic

AUC Area Under the Curve

pypi Python Packaging Index

List of Listings

1	generator network	29
2	discriminator network	31
3	training loop	33
4	training loop	37
5	training loop	42
6	gridsearch simplified	43
7	base model	69

List of Figures

1	Decision Map for choosing the right estimator, from scikit learn https://scikit-learn.org/stable/tutorial/machine_learning_map/index.html	3
2	from El Emam et al. (2020)	6
3	from Patki et al. (2016)	8
4	from Ping et al. (2017)	10
5	showing the decision boundary of an SVM classifier trained on data generated on original data (left), data generated using the DS (middle) and SDV (right) from Hittmeir et al. (2019)	11
6	from Chawla et al. (2002)	13
7	showing classical image transformations from Perez and Wang (2017)	15
8	from Goodfellow et al. (2014)	21
9	Initial simple dense GAN - left side shows the losses of generator and discriminator, right side shows the probabilities assigned to real and fake samples by the discriminator	24

10	Dense GAN, 3 layers, 64 neurons/layer; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator	25
11	Dense GAN, 3 layers, 128 neurons/layer, reduced learning rate and dropout in discriminator - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator .	26
12	Architecture of Discriminator and Generator	27
13	Dense GAN, 2 layers, 32 neurons/layer; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator	29
14	EM / Wasserstein distance formula	35
15	Wasserstein GAN, 3 layers; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator	36
16	Wasserstein GAN with gradient penalty, 2 layers; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator	38
17	Experiment with data sizes 30%, 40% of original data	46
18	Experiment with data sizes 50%, 60% of original data	46
19	Experiment with data sizes 70%, 80% of original data	47
20	Experiment with data size 100% of original data	48
21	Experiment with data sizes +300% - +2000% original data	49
22	Comparing all model type performances, training data size boosted by 20% with Wasserstein-GP generator	56

23	Wasserstein GAN with gradient penalty, 2 layers; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator	58
24	Comparing all model type performances, training data size boosted by 20% with Wasserstein-GP generator	59
25	Wasserstein GAN with gradient penalty, 2 layers; left - losses of generator/discriminator right - the probabilities real/fake assigned by the discriminator	63
26	Comparing all model type performances, training data size boosted by 20% with Wasserstein-GP generator	64
27	Comparing all model type performances, training data entirely replaced by synthetic data of same size with Wasserstein-GP generator	66
28	Comparing all model type performances, training data entirely replaced by synthetic data of same size with Wasserstein-GP generator	67
29	Comparing all model type performances, training data entirely replaced by synthetic data of same size with Wasserstein-GP generator	68
30	model accuracy history during training; data entirely replaced by synthetic data of same size with Wasserstein-GP generator . .	71